

Parallel Gradient Descent Algorithm with *MPI* on *Distributed Architecture*

Dr. Süha Tuna

FSMVU – Department of Computer Engineering

Content

1. Problem Definition
2. Gradient Descent Algorithm
3. Thinking Parallel
4. Implementation Details


Problem Definition

- ▶ Solving an **underdetermined** linear equation system
- ▶ $Ax = b$
- ▶ A and b is given, x is unknown
- ▶ A is size of $n_1 \times n_2$ where $n_2 > n_1$
 - ▶ # of unknowns > # of equations
- ▶ No unique solution!
- ▶ Conventional iterative solvers may not work!
 - ▶ Rectangularity

Gradient Descent Algorithm

- ▶ Minimize $E(\mathbf{x}) = \frac{1}{2} ||A \mathbf{x} - \mathbf{b}||_2^2$
 - ▶ Convex cost function
- ▶ $\nabla_{\mathbf{x}} E(\mathbf{x}) = A^T (A \mathbf{x} - \mathbf{b})$
- ▶ Initial guess: \mathbf{x}_0
- ▶ Learning rate: α
- ▶ GD iterations: $\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla_{\mathbf{x}} E(\mathbf{x}_k), \quad k = 0, 1, 2, \dots$
- ▶ Converges to global minimum, if
 - ▶ Cost is convex
- ▶ α should be examined and chosen carefully!

Thinking Parallel



Parallel Implementation - I


► Main iteration:

$$\text{► } \mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \mathbf{A}^T (\mathbf{A} \mathbf{x}_k - \mathbf{b})$$

1. Matrix vector multiplication: $\mathbf{A} \mathbf{x} = \mathbf{r}$

1. Start with an initial guess: \mathbf{x}_0
2. MPI_Scatter **rows** of \mathbf{A} matrix to each process' \mathbf{A}_{local}
 1. Can use derived data type: MPI_Type_contiguous
3. MPI_Broadcast \mathbf{x} to each process
4. Perform $\mathbf{A}_{local} \times \mathbf{x} = \mathbf{r}_{local}$
5. No need to MPI_Gather for \mathbf{r}_{local}


Thinking Parallel



Parallel Implementation - II

1. Vector subtraction: $\mathbf{r} - \mathbf{b} = \mathbf{y}$
 1. \mathbf{r}_{local} is already calculated
 2. MPI_Scatter \mathbf{b} vector to each process' \mathbf{b}_{local}
 3. Perform local vector subtraction $\mathbf{r}_{local} - \mathbf{b}_{local} = \mathbf{y}_{local}$


Thinking Parallel



Parallel Implementation - III

1. Matrix vector multiplication: $A^T \mathbf{y} = \mathbf{u}$
 1. MPI_Scatter **columns** of A matrix to each process' A_{local}^T
 1. Can use derived data type: MPI_Type_vector
 2. Do not forget to use MPI_Type_create_resized
 2. MPI_Allgather \mathbf{y}_{local} 's from each process to \mathbf{y} vector
 3. Perform $A_{local}^T \times \mathbf{y} = \mathbf{u}_{local}$
 4. No need to MPI_Gather for \mathbf{u}_{local}

Thinking Parallel



Parallel Implementation - IV

1. Scalar multiplication: αu_{local}
2. MPI_Scatter x to each process' x_{local}
3. Vector subtraction: $x := x - \alpha u$
 1. Perform local subtraction: $x_{local} := x_{local} - \alpha u_{local}$
4. Repeat steps 1 to 4 till convergence or the last iteration

Implement the algorithm using MPI in C.