



Shared Memory Programming with OpenMP

Dr. Süha Tuna

*Fatih Sultan Mehmet Vakıf University,
Computer Engineering*

2017-2018 Spring

Outline - I

- Shared Memory Systems
 - Threaded Programming Model
 - Thread Communication
 - Synchronisation
 - Parallel Loops
 - Reductions
- OpenMP Fundamentals
 - Basic Concepts in OpenMP
 - History of OpenMP
 - Compiling and Running OpenMP Programs

Outline - II


- Parallel Regions
 - Parallel region directive
 - Some useful functions
 - Shared and Private variables
 - Reductions
- Work sharing
 - Parallel for/DO loops
 - Scheduling for loops
 - Single directive
 - Master directive
- Synchronisation

Shared Memory Systems

Shared Memory Systems

- Threaded programming is most often used on shared memory parallel computers.
- A shared memory computer consists of a number of processing units (CPUs) together with some memory.
- Key feature of shared memory systems is *single address space* across the whole memory system.
 - every CPU can read or write all memory locations in the system
 - one logical memory space
 - all CPUs refer to a memory location using the same address


Conceptual Model



Real Hardware

- Real shared memory hardware is more complicated than this ...
 - Memory may be split into multiple smaller units
 - There may be multiple levels of cache memory
 - some of these levels may be shared between subsets of processors
 - The interconnect may have a more complex topology
- ... but a single space address is still supported
 - Hardware complexity can affect the performance of programs, but not their correctness.

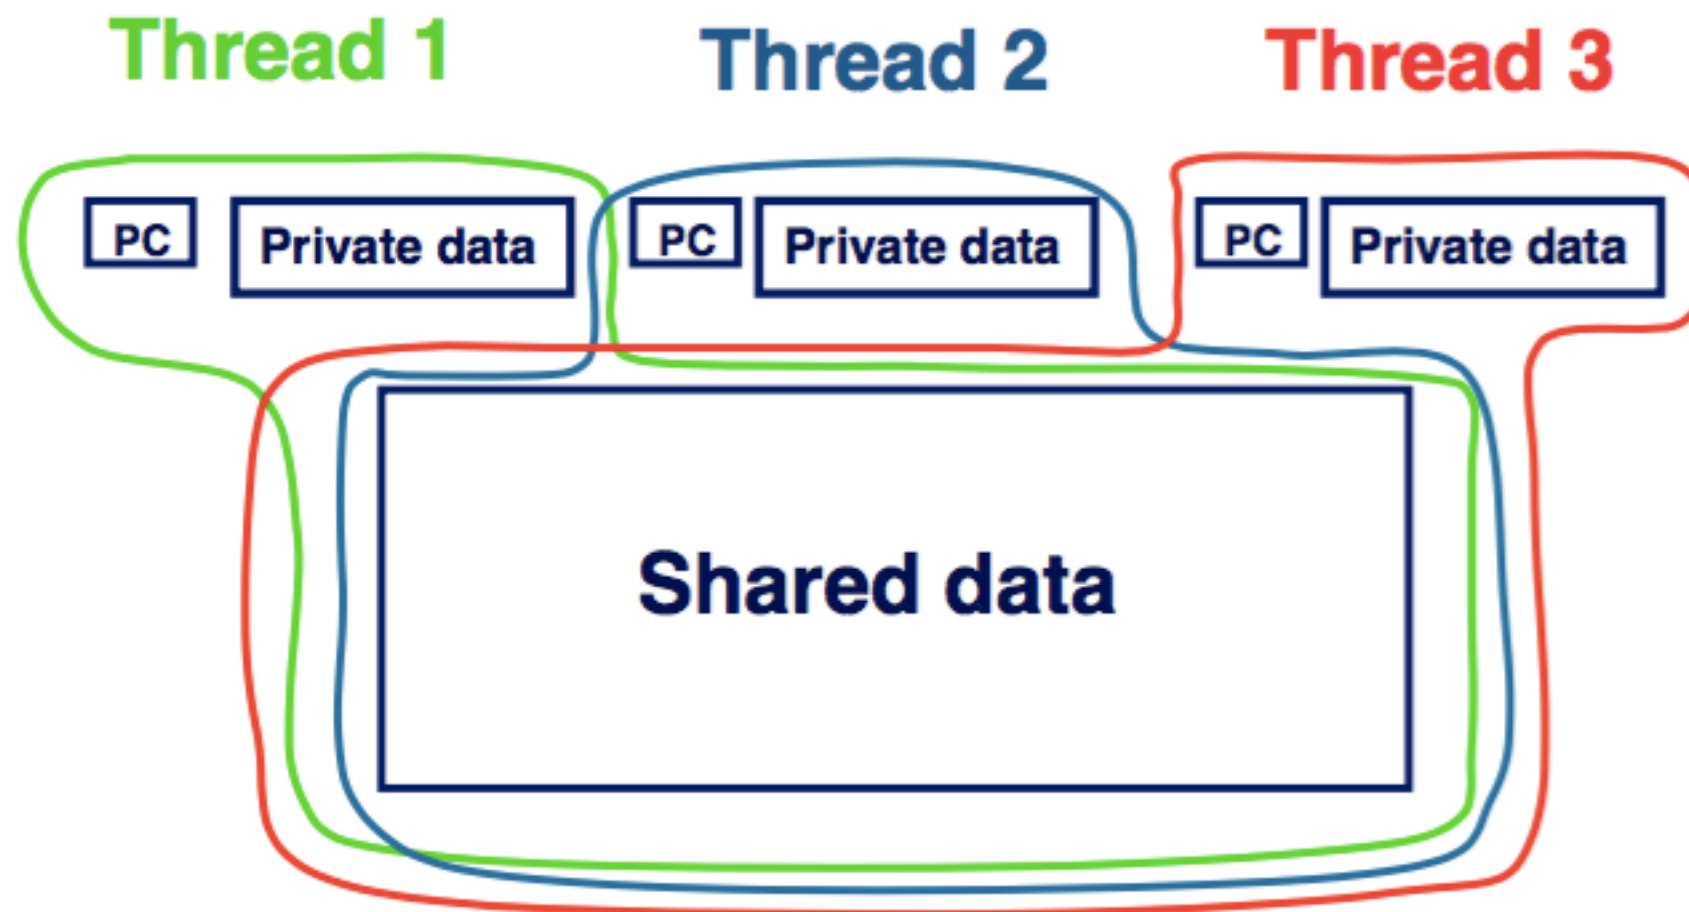
Real Hardware Example



Threaded Programming Model

- The programming model for shared memory is based on the notions of threads
 - threads are like processes, except that threads can share memory with each other (as well as having private memory)
- Shared data can be accessed by all threads
- Private data can only be accessed by the owning thread
- Different threads can follow different flows of control through the same program
 - each thread has its own program counter
- Usually run one thread per CPU/core
 - but could be more
 - can have hardware support for multiple threads per core


Threads (cont.)



Thread Communication

- In order to have useful parallel programs, threads must be able to exchange data with each other
- Threads communicate with each via reading and writing shared data
 - thread 1 writes a value to a shared variable A
 - thread 2 can then read the value from A
- Note: there is no notion of messages in this model


Thread Communication



Synchronisation

- By default, threads execute **asynchronously**
- Each thread proceeds through program instructions independently of other threads
- This means we need to ensure that actions on shared variables occur in the correct order: e.g.
 - thread 1 must write variable A before thread 2 reads it
 - or
 - thread 1 must read variable A before thread 2 writes it
- Note that updates to shared variables (e.g. **$a = a + 1$**) are not atomic!
- If two threads try to do at the same time, one of the updates may get overwritten.

Synchronisation Example



Parallel Loops

- Loops are the main source of parallelism in many applications
- If the iterations of a loop are independent (can be done any order) then we can share out the iterations between different threads
- e.g. if we have two threads and the loop

```
for (i=0; i<100; i++) {  
    a[i] += b[i];  
}
```

 - we could do iteration 0-49 on one thread and iterations 50-99 on the other.

for ($i=0; i<100; i++$)
 $a[i] += b[i];$

Parallelizable

for ($j=0; j<100; j++$)
 $a[j] += 2;$

} parallelizable

for ($j=0; j<100; j++$)
 $a[j] = b[j] + b[j+1];$

Parallelizable

$$\begin{aligned} a_0 &= b_0 + b_1^{(4)} \\ a_1 &= b_1 + b_2 \\ a_2 &= b_2 + b_3^{(5)} \\ &\vdots \end{aligned}$$

$$\begin{aligned} a_{26} &= b_{26} + b_{27}^{(3)} \\ &\vdots \\ a_{93} &= b_{93} + b_{94}^{(1)} \\ &\vdots \\ a_{99} &= b_{99} + b_{100}^{(2)} \end{aligned}$$

for ($k=0; k<100; k++$)
 $a[k+2] = a[k+1] + a[k];$

$$① a_2 = a_1 + a_0$$

$$② a_3 = a_2 + a_1$$

$$③ a_4 = a_3 + a_2$$

\vdots


①

} order matters ✓
 loop is not independent,
 thus un-parallelizable!

Reductions

- A *reduction* produces a single value from associative operations such as addition, multiplication, max, min, and, or.

- For example:



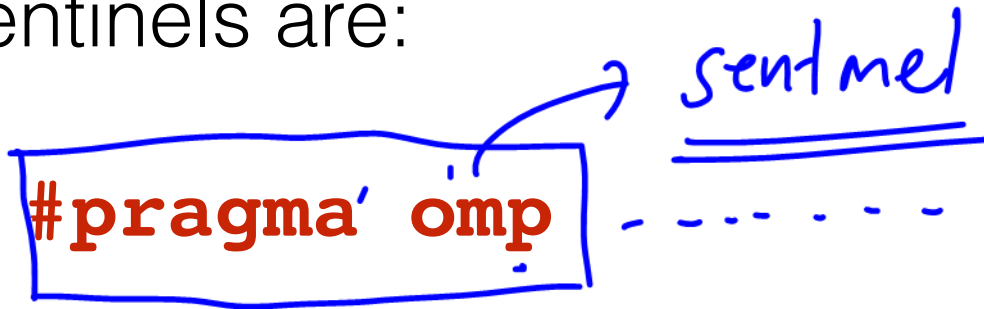
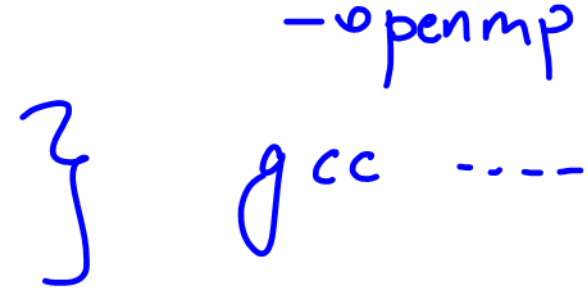
- Allowing only one thread at a time to update **b** would remove all parallelism
- Instead, each thread can accumulate its own private copy, then these copies are reduced to give final result
- If the number of operations is much larger than the number of threads, most of the operations can proceed in parallel

OpenMP Fundamentals

What is OpenMP

- OpenMP is an API designed for programming shared memory parallel computers
- OpenMP uses the concepts of threads
- OpenMP is a set of extensions to C, C++ and Fortran
- The extensions consist of:
 - Compiler directives → MPI-Send, MPI-Recv, +, -, =
 - Runtime library routines → size - np (8)
MPI-Comm-size, MPI-Comm-rank
 - Environment variables → n threads → BASH
OMP_NUM_THREADS


Directives and Sentinels

- A directive is a special line of source code with meaning only to certain compilers
- A directive is distinguished by a sentinel at the start of the line
- OpenMP sentinels are:
 - C / C++: **#pragma omp**  
 - Fortran : **! \$OMP**
- This means that OpenMP directives are ignored if the code is compiled as regular sequential C/C++/Fortran

Parallel Region

- The *parallel region* is the basic parallel construct in OpenMP
- A parallel region defines a section of a program
- Program begins execution on a single thread (the master thread)
- When the first parallel region is encountered, master thread creates a team of threads (fork/join model)
- Every thread executes the statements which are inside the parallel region
- At the end of the parallel region, the master thread waits for the other threads to finish, and continues executing the next statements

Parallel Region



^(Default) Shared and Private Data

- Inside a parallel region, variables can either be *shared* or *private*
first private *last private*
- All threads see the same copy of shared variables
- All threads can read or write shared variables
- Each thread has its own copy of private variables: these are invisible to other threads
- A private variable can only be read or written by its own thread

Parallel Loops

work

- In a parallel region, all threads execute the same code
- OpenMP has also directives which indicate that work should be divided up between threads, not replicated
 - this is called worksharing
- Since loops are the main source of parallelism in many applications, OpenMP has an extensive support for parallelising loops
- There are a number of options to control which loop iterations are executed by which threads
- It is up to programmer to ensure that the iterations of a parallel loop are *independent*
- Only loops where the iteration count can be computed before the execution of the loop begins can be parallelised in this way

Synchronisation

- The main synchronisation concepts used in OpenMP are:

- Barrier


- all threads must arrive at a barrier before any thread can proceed past it
- e.g. delimiting phases of computation

- Critical regions

- a section of code which only one thread at a time can enter
- e.g. modification of shared variables

- Atomic update

- an update to a variable which can be performed only by one thread at a time
- e.g. modification of shared variables (special case)



$(\text{shared}) \text{ count}++;$ } only one single var can be updated!

Brief History of OpenMP

- Historical lack of standardisation in shared memory directives
 - each hardware vendor provided a different API
 - mainly directive based
 - almost all for Fortran
 - hard to write portable code
- OpenMP forum is set up by Digital, IBM, Intel, KAI and SGI. Now includes most major vendors (and some academic organisations)
- OpenMP Fortran standard released in October 1997, minor revision (1.1) in November 1999, Major revision (2.0) in November 2000
- OpenMP C/C++ standard released October 1998. Major revision (2.0) in March 2002

History (cont.)

- Combined OpenMP C/C++/Fortran standard (2.5) released in May 2005
 - no new features, but extensive rewriting and clarification
- Version 3.0 released in May 2008
 - new features, including tasks, better support for loop parallelism and nested parallelism
- Version 3.1 released in June 2011
 - corrections and some minor new features
 - most current compilers support this
- Version 4.0 released in July 2013
 - accelerator offloading, thread affinity, more task support
 - now appearing in implementations
- Version 4.5 released in November 2015
 - corrections and a few new features



Compiling and Running OpenMP Programs

- OpenMP is built-in to most of the compilers you are likely to use
- To compile OpenMP program you need to add a (compiler-specific) flag to your compile and link commands
 - **-fopenmp** for gcc/gfortran
 - **-openmp** for Intel compilers
- The number of threads which will be used is determined at runtime by **OMP_NUM_THREADS** environment variable
 - set this before you run the program
 - e.g. **export OMP_NUM_THREADS=4**
- Run in the same way you would a sequential program
 - type the name of the executable

Exercise

- “Hello World” program
- Aim: to compile and run a trivial OpenMP program
- Vary the number of threads using the **OMP_NUM_THREADS** environment variable
- Run the code several times. Is the output always the same?

Exercise II.

PI Calculation with OpenMP

$$\int_0^1 \frac{4}{1+x^2} dx = 4 \arctan x \Big|_0^1 = 4 \left(\underbrace{\arctan 1}_{\frac{\pi}{4}} - \underbrace{\arctan 0}_0 \right) = \pi$$




N steps

$$h = \frac{1-0}{N} = \frac{1}{N} \Rightarrow Nh = 1$$

$$h \cdot \frac{4}{1+h^2} + h \cdot \frac{4}{1+(2h)^2} + h \cdot \frac{4}{1+(3h)^2} + \dots + h \cdot \frac{4}{1+(nh)^2}$$

$$\approx \pi$$

Exercise: Monte - Carlo Simulation



$$x^2 + y^2 = 1$$

if $x^2 + y^2 \leq 1$

$(x, y) \in A \rightarrow \text{counter}++;$

else

$(x, y) \notin A$

$$x^2 + y^2 = 1$$

$$N = 100K$$

$$\left. \begin{array}{l} 0 \leq x \leq 1 \\ 0 \leq y \leq 1 \end{array} \right\} \text{Random}$$

for $i = 1 : N$

$$\frac{3.141590...}{4} \approx \underline{\underline{0.79...}}$$

$$B = N - A$$

$$\frac{A}{A+B} \approx \frac{\pi/4}{1}$$

$$N = \frac{A}{A+B} \rightarrow 1$$

P_0
P_1
P_2
P_3


1D Cartesian Topo.

↳ line top

$$\varepsilon_0 = \|A x_0 - b\|_2 \rightarrow$$

$$\varepsilon_i = \|A x_i - b\|_2$$

$$\frac{\varepsilon_i}{\varepsilon_0} \rightarrow \frac{50}{100} = 0.5$$



Parallel Regions

Parallel Region Directive


- Code within a parallel region is executed by all threads
- Syntax:
 - C/C++:

```
#pragma omp parallel  
{  
    block  
}
```
 - Fortran:

```
!$OMP PARALLEL  
    block  
!$OMP END PARALLEL
```

Parallel Region Directive (cont.)

```
fred();  
#pragma omp parallel  
{  
    billy();  
}  
daisy();
```



Useful Functions

- Often useful to find out number of threads being used

- Fortran:

```
USE OMP_LIB  
INTEGER FUNCTION OMP_GET_NUM_THREADS()
```

- C/C++:

```
#include <omp.h>  
int omp_get_num_threads(void);
```

- Note: returns 1 if called outside parallel region!

Useful Functions (cont.)

- Also useful to find out number of the executing thread

- Fortran:

```
USE OMP_LIB  
INTEGER FUNCTION OMP_GET_THREAD_NUM( )
```

- C/C++:

```
#include <omp.h>  
int omp_get_thread_num(void);
```

- Note: Takes value between 0 and

```
OMP_GET_NUM_THREADS( ) - 1
```

Clauses

- Specify additional information in the parallel region directive through *clauses*:
 - C/C++: **#pragma omp parallel [*clauses*]**
 - Fortran: **!\$OMP PARALLEL [*clauses*]**
- Clauses are comma or space separated in Fortran, space separated in C/C++

Shared and Private Variables

- Inside a parallel region, variables can be either **shared** (all threads see same copy) or **private** (each thread has its own copy)
- **shared**, **private** and **default** are OpenMP clauses
 - C/C++: **shared(list)**
private(list)
default(shared|none)
 - Fortran: **SHARED(list)**
PRIVATE(list)
DEFAULT(SHARED|PRIVATE|NONE)


Shared and Private (cont.)

- On entry to a parallel region, private variables are uninitialised
- Variables declared inside the scope of the parallel region are automatically private
- After the parallel region ends, the original variable is unaffected by any changes to private copies
- Not specifying a **DEFAULT** clause is the same as specifying **DEFAULT (SHARED)**
 - Danger!
 - Always use **DEFAULT (NONE)**

Shared and Private (cont.)

- Example: each thread initializes its own column of a shared array

```
!$OMP PARALLEL  DEFAULT (NONE), PRIVATE (I, MYID),  
!$OMP&  SHARED(A,N)  
    MYID = OMP_GET_THREAD_NUM() + 1  
    DO I = 1, N  
        A(I, MYID) = 1.0  
    END DO  
!$OMP END PARALLEL
```



Multi-line Directives

C/C++:

```
#pragma omp parallel default(none) \  
private(i,myid) shared(a,n)
```

Fortran: fixed source form

```
!$OMP PARALLEL DEFAULT(NONE), PRIVATE(I,MYID),  
!$OMP& SHARED(A,N)
```

Fortran: free source form

```
!$OMP PARALLEL DEFAULT(NONE), PRIVATE(I,MYID), &  
!$OMP SHARED(A,N)
```

Initializing Private Variables

- Private variables are uninitialized at the start of the parallel region
- If we wish to initialize them, we use **FIRSTPRIVATE** clause:
 - C/C++: **firstprivate(*list*)**
 - Fortran: **FIRSTPRIVATE(*list*)**

Initializing Private Variables (cont.)

```
b = 23.0;
```

```
. . . . .
```

```
#pragma omp parallel firstprivate(b),  
private(i, myid)
```

```
{
```

```
    myid = omp_get_thread_num();
```

```
    for (i=0; i<n; i++) {
```

```
        b += c[myid][i]; → Each thread sums its  
row up to b
```

```
    }
```

```
    c[myid][n] = b;
```

```
}
```

	0	1	...	N-1	N
T ₀	0	0	...	0	X
T ₁	0	0	...	0	X
T ₂	0	0	...	0	X

23 + ...
No race
condition!

Reductions


- A *reduction* produces a single value from associative operations such as addition, multiplication, max, min, and, or
- Would like each thread to reduce into a private copy, then reduce all these to give final result

- Use **REDUCTION** clause:

4 threads

- C/C++: **reduction(op: list)**

- Fortran: **REDUCTION(op: list)**



- Can have reduction arrays in Fortran, but not in C/C++

Reductions (cont.)

Value in original variable is saved

Each thread gets a private copy of **b**, initialized to 0

B = 10

!\$OMP PARALLEL REDUCTION (+:B),

!\$OMP& PRIVATE(I, MYID)

MYID = OMP_GET_THREAD_NUM() + 1

DO I = 1, N

B = B + C[I][MYID]

All accesses inside the parallel region are to the private copies


END DO

!\$OMP END PARALLEL

At the end of the parallel region, all the private copies are added into the original variable

A = B

Exercise



- Area of the Mandelbrot set
- Aim: introduction to using parallel regions
- Estimate the area of the Mandelbrot set by Monte Carlo sampling
 - Generate a grid of complex numbers in a box surrounding the set
 - Test each number to see if it is in the set or not
 - Ratio of points inside a total number of points gives an estimate of the area
 - Testing of points is independent - parallels with a parallel region

Worksharing

Worksharing Directives

- Directives which appear inside a parallel region and indicate how work should be shared out between threads are
 - Parallel DO/for loops
 - Single directive
 - Master directive

Parallel DO/for Loops

- Loops are the most common source of parallelism in most codes. Therefore, parallel loop directives are very important!
- A parallel DO/for loop divides up the iterations of the loop between threads
- The loop directive appears inside a parallel region and indicates that the work should be shared out between threads, instead of replication
- There is a synchronisation point at the end of the loop: all threads must finish their iterations before any thread can proceed

Parallel DO/for Loops (cont.)

Syntax:

Fortran: **!\$OMP DO** *[clauses]*
 DO loop
 !\$OMP END DO

C/C++: **#pragma omp for** *[clauses]*
 for loop

Restrictions in C/C++

- Because the for loop in C is a general while loop, there are restrictions on the form it can take
- It has two determinable trip count - it must be of the form
 - **for (var = a; var *logical-op* b; incr-exp)**
 - where *logical-op* is one of <, <=, >, >=
 - and *incr-exp* is **var = var +/- incr** or semantic equivalent such as **var++**
 - also can not modify **var** within the loop body

Parallel Loops (Example)

```
! $OMP PARALLEL
! $OMP DO
    DO i=1,n
        b(i)=(a(i)-a(i-1))*0.5
    END DO
! $OMP END DO
! $OMP END PARALLEL
```

```
#pragma omp parallel
{
    pragma omp for
    #pragma omp for
    for (int i=1, i<=n, i++) {
        b(i)=(a(i)-a(i-1))*0.5;
    }
}
```

Parallel DO/for Directive

- This construct is common that there is shorthand form which combines parallel region and DO/for loops
 - C/C++: **#pragma omp parallel for [*clauses*]**
for loop
 - Fortran: **!\$OMP PARALLEL DO [*clauses*]**
do loop
!\$OMP END PARALLEL DO

Clauses

- DO/for directive can take PRIVATE, FIRSTPRIVATE and REDUCTION clauses which refer to the scope of the loop
- Note that the parallel loop variable is PRIVATE by default
 - loop indices are private by default in Fortran, but not in C
- PARALLEL DO/for directive can take all clauses available for PARALLEL directive
- PARALLEL DO/for is not the same as DO/for or the same as PARALLEL

Parallel DO/for Loops (cont.)

- With no additional clauses, the DO/for directive will partition the iterations as equally as possible between the threads
- However this is implementation dependent, and there is still some ambiguity
 - e.g. 7 iterations, 3 threads. Could partition as 3+3+1 or 3+2+2

SCHEDULE Clause

- The SCHEDULE clause gives a variety of options for specifying which loop iteration are executed by which thread
- Syntax:
 - C/C++: **`schedule(kind[, chunksize])`**
 - Fortran: **`SCHEDULE(kind[, chunksize])`**
 - where kind is one of
 - **STATIC, DYNAMIC, GUIDED, AUTO or RUNTIME**
 - and chunksize is an integer expression with positive value
- e.g. **`!$OMP DO SCHEDULE(DYNAMIC, 4)`**


STATIC Schedule

- With no *chunksize* specified, the iteration space is divided into (approximately) equal chunks, and one chunk is assigned to each thread in order (**block** schedule)
- If *chunksize* is specified, the iteration space is divided into chunks, each of *chunksize* iterations, and the chunks are assigned cyclically to each thread in order (**block cyclic** schedule)

STATIC Schedule



SCHEDULE (STATIC)



SCHEDULE (STATIC, 4)

DYNAMIC Schedule

- DYNAMIC schedule divides the iteration space up into chunks of size *chunksize*, and assigns them to threads on a first-come-first-served basis
- i.e. as a thread finish a chunk, it is assigned the next chunk in the list
- When no *chunksize* is specified, it defaults to 1

150


4 threads

15	15	15	15	→ 60
8	8	8	8	→ 32
4	4	4	4	→
2	2	2	2	→
1	1	1	1	→
1	1	1	1	→

GUIDED Schedule

- GUIDED schedule is similar to DYNAMIC, but the chunk starts off large and gets smaller exponentially
- The size of the next chunk is proportional to the number of remaining iteration divided by the number of threads
- The *chunksize* specifies the minimum size of the chunks
- When no *chunksize* is specified, it defaults to 1

DYNAMIC and GUIDED Schedules



1

SCHEDULE (DYNAMIC, 3)

46



1

SCHEDULE (GUIDED, 3)

46

AUTO Schedule

- Lets the runtime have full of freedom to choose its own assignment of iterations to threads
- If the parallel loop is executed many times, the runtime can evolve a good schedule which has good load balance and low overheads

Choosing a Schedule

- STATIC best for load balanced loops - least overhead
- STATIC, n good for loops with mild or smooth load imbalance, but can induce overheads
- DYNAMIC useful if iterations have widely varying loads, but ruins data locality
- GUIDED often less expensive than DYNAMIC, but beware of loops where the first iterations are the most expensive
- AUTO may be useful if the loop is executed many times over

SINGLE Directive

- Indicates that a block of code is to be executed by a **single thread only**
- The first thread to reach the SINGLE directive will execute the block
- There is a **synchronisation** point at the end of the block: all other threads wait until block has been executed

SINGLE Directive (cont.)


Syntax:

Fortran **!\$OMP SINGLE [*clauses*]**
 block
 !\$OMP END SINGLE

C/C++: **#pragma omp single [*clauses*]**
 structured block

SINGLE Directive (cont.)

```
#pragma omp parallel
{
    setup(x);
    #pragma omp single
    {
        input(y);
    }
    work(x,y);
}
```



SINGLE Directive (cont.)

- SINGLE directive can take PRIVATE and FIRSTPRIVATE clauses
- Directive must contain a structured block: can not branch into or out of it

MASTER Directive

- Indicates that a block of code should be executed by the **master thread** (thread 0) **only**
- There is **no synchronisation** at the end of the block: other threads skip the block and continue executing

MASTER Directive (cont.)

Syntax:

Fortran **!\$OMP MASTER**
 block
 !\$OMP END MASTER

C/C++: **#pragma omp master**
 structured block

Synchronisation

Why is It Required?

- Need to synchronise actions on shared variables
- Need to ensure correct ordering of reads and writes
- Need to protect updates to shared variables (not atomic by default)


BARRIER Directive

- No thread can proceed reached a barrier until all the other threads have arrived
- Note that there is an implicit barrier at the end of DO/for, SECTIONS and SINGLE directives
- Syntax:
 - C/C++: **#pragma omp barrier**
 - Fortran: **!\$OMP BARRIER**
- Either all threads or none must encounter the barrier: otherwise DEADLOCK!

BARRIER Directive (cont.)

Example:

```
!$OMP PARALLEL PRIVATE(I,MYID,NEIGHB)
  myid = omp_get_thread_num()
  neighb = myid - 1
  if (myid.eq.0) neighb = omp_get_num_threads()-1
  ...
  a(myid) = a(myid)*3.5
!$OMP BARRIER
  b(myid) = a(neighb) + c
  ...
!$OMP END PARALLEL
```



- Barrier required to force synchronisation on **a**

Critical Sections

- A critical section is a block of code which can be executed by only one thread at a time
- Can be used to protect updates to shared variables
- The CRITICAL directive allows critical sections to be named
- If one thread is in a critical section with a given name, no other thread may be in a critical section with the same name (though they can be in critical sections with other names)

Critical Directive

Fortran **!\$OMP CRITICAL [(*name*)]**
 block
 !\$OMP END CRITICAL [(*name*)]

C/C++: **#pragma omp critical [(*name*)]**
 structured block

- In Fortran, the names on the directive pair must match
- If the name is omitted, a null name is assumed (all unnamed critical sections effectively have the same null name)

Critical Directive (cont.)

Example: Pushing and popping a task stack

```
!$OMP PARALLEL SHARED (STACK) , PRIVATE (INEXT , INEW)  
    ...  
!$OMP CRITICAL (STACKPROT)  
    inext = getnext(stack)  
!$OMP END CRITICAL (STACKPROT)  
    call work(inext,inew)  
!$OMP CRITICAL (STACKPROT)  
    if (inew .gt. 0) call putnew(inew,stack)  
!$OMP END CRITICAL (STACKPROT)  
    ...  
!$OMP END PARALLEL
```

Atomic Directive

- Used to protect a single update to a shared variable.
- Applies only to a single statement.
- Syntax:

Fortran: **!\$OMP ATOMIC**
statement

where *statement* must have one of these forms:

$x = x \text{ op } \text{expr}, \quad x = \text{expr op } x, \quad x = \text{intr} (x, \text{expr}) \text{ or}$
 $x = \text{intr}(\text{expr}, x)$

op is one of +, *, -, /, .and., .or., .eqv., or .neqv.

intr is one of **MAX**, **MIN**, **IAND**, **IOR** or **IEOR**

Atomic Directive (cont.)

C/C++: `#pragma omp atomic`
statement

where *statement* must have one of the forms:

x binop = expr, *x++*, *++x*, *x--*, or *--x*


and *binop* is one of *+*, ***, *-*, */*, *&*, *^*, *<<*, or *>>*

- Note that the evaluation of *expr* is not atomic.
- May be more efficient than using CRITICAL directives, e.g. if different array elements can be protected separately.
- No interaction with CRITICAL directives

Atomic Directive (cont.)

Example (compute degree of each vertex in a graph):

```
#pragma omp parallel for
    for (j=0; j<nedges; j++){
#pragma omp atomic
        degree[edge[j].vertex1]++;
#pragma omp atomic
        degree[edge[j].vertex2]++;
    }
```



QUESTIONS
or
COMMENTS!